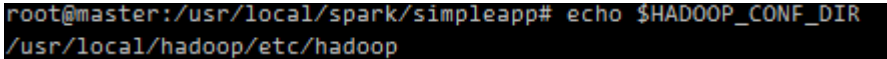


Spark Tutorial

Jinlai Xu

Preparation

- Check HDFS and YARN services and environment configuration
 - jps
 - echo \$HADOOP_CONF_DIR 
- Download Spark package
 - wget <http://apache.mirrors.pair.com/spark/spark-2.3.0/spark-2.3.0-bin-hadoop2.7.tgz>
- Unpack
 - tar zxvf [spark-2.3.0-bin-hadoop2.7.tgz](#)
- Try Spark shell (local mode)
 - ./bin/spark-shell --master local[2]

```
val NUM_SAMPLES = 1000
val count = sc.parallelize(1 to NUM_SAMPLES).filter { _ =>
  val x = math.random
  val y = math.random
  x*x + y*y < 1
}.count()
println(s"Pi is roughly ${4.0 * count / NUM_SAMPLES}")
```

Run Spark Shell with YARN

- `./bin/spark-shell --master yarn --deploy-mode client`
 `val NUM_SAMPLES = 1000`
 `val count = sc.parallelize(1 to NUM_SAMPLES).filter { _ =>`
 `val x = math.random`
 `val y = math.random`
 `x*x + y*y < 1`
 `}.count()`
 `println(s"Pi is roughly ${4.0 * count / NUM_SAMPLES}")`

Run Spark example program with YARN

- `./bin/spark-submit --class org.apache.spark.examples.SparkPi \
--master yarn \
--deploy-mode cluster \
--driver-memory 512m \
--executor-memory 512m \
--executor-cores 1 \
--queue default \
examples/jars/spark-examples*.jar \
10`

Package and Run your Spark JAVA program

- Directory Structure:
 - pom.xml
 - src/main/java/your_program.java
- pom.xml
 - <project>
 - <groupId>edu.berkeley</groupId>
 - <artifactId>simple-project</artifactId>
 - <modelVersion>4.0.0</modelVersion>
 - <name>Simple Project</name>
 - <packaging>jar</packaging>
 - <version>1.0</version>
 - <dependencies>
 - <dependency> <!-- Spark dependency -->
 - <groupId>org.apache.spark</groupId>
 - <artifactId>spark-core_2.11</artifactId>
 - <version>2.3.0</version>
 - </dependency>
 - </dependencies>
 - </project>

Package and Run your Spark JAVA program

- SimpleApp.java

- `/* SimpleApp.java */`
- `import org.apache.spark.api.java.*;`
- `import org.apache.spark.SparkConf;`
- `import org.apache.spark.api.java.function.Function;`
- `public class SimpleApp {`
- `public static void main(String[] args) {`
- `String logFile = "YOUR_SPARK_HOME/README.md"; // Should be some file on your system`
- `SparkConf conf = new SparkConf().setAppName("Simple Application");`
- `JavaSparkContext sc = new JavaSparkContext(conf);`
- `JavaRDD<String> logData = sc.textFile(logFile).cache();`
- `long numAs = logData.filter(new Function<String, Boolean>() {`
- `public Boolean call(String s) { return s.contains("a"); }`
- `}).count();`
- `long numBs = logData.filter(new Function<String, Boolean>() {`
- `public Boolean call(String s) { return s.contains("b"); }`
- `}).count();`
- `System.out.println("Lines with a: " + numAs + ", lines with b: " + numBs);`
- `sc.stop();`
- `}`
- `}`

Package and Run your Spark JAVA program

- Use maven to package the program:
 - Make sure you install maven on the client:
 - `sudo apt-get install maven`
 - Package the program:
 - `mvn package`

```
[INFO] Building jar: /usr/local/spark-2.1.0-bin-hadoop2.7/simpleapp/target/simpl
e-project-1.0.jar
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 20.439 s
[INFO] Finished at: 2017-03-27T14:54:34+00:00
[INFO] Final Memory: 35M/190M
[INFO] -----
```

Package and Run your Spark JAVA program

- Directory Structure:
 - pom.xml
 - src/main/java/your_program.java
 - **target/your_program*.jar**
- Run the program with spark-submit
 - `../bin/spark-submit --class "SimpleApp" \`
`--master yarn \`
`--deploy-mode cluster \`
`--driver-memory 1g \`
`--executor-memory 1g \`
`--executor-cores 1 \`
`--queue default \`
`target/simple*.jar`

```
17/03/27 15:04:53 INFO yarn.Client: Application report for application_1490620408521_0007 (state: RUNNING)
17/03/27 15:04:54 INFO yarn.Client: Application report for application_1490620408521_0007 (state: RUNNING)
17/03/27 15:04:55 INFO yarn.Client: Application report for application_1490620408521_0007 (state: FINISHED)
17/03/27 15:04:55 INFO yarn.Client:
    client token: N/A
    diagnostics: N/A
    ApplicationMaster host: 162.243.40.202
    ApplicationMaster RPC port: 0
    queue: default
    start time: 1490627066999
    final status: SUCCEEDED
    tracking URL: http://master:8088/proxy/application_1490620408521_0007/
    user: root
17/03/27 15:04:55 INFO util.ShutdownHookManager: Shutdown hook called
17/03/27 15:04:55 INFO util.ShutdownHookManager: Deleting directory /tmp/spark-1bda6d45-37ff-4834-ac23-a1c4db31cc5c
```


Q & A

FAQ

- The compatibility of Spark, Hadoop and JAVA
 - From Spark Official website:
 - Spark runs on Java 7+, Python 2.6+/3.4+ and R 3.1+. For the Scala API, Spark 2.1.0 uses Scala 2.11. You will need to use a compatible Scala version (2.11.x).
 - Note that support for Java 7 and Python 2.6 are **deprecated (do not mean unsupported)** as of Spark 2.0.0, and support for Scala 2.10 and versions of Hadoop before 2.6 are **deprecated** as of Spark 2.1.0, and may be removed in Spark 2.2.0.
- The comparison in Part 3
 - Required: compare the performance of using the cached RDD feature and without using the cached RDD feature.
 - Alternative: compare the performance of above two with the same workload of using Hadoop MapReduce

FAQ

- Do I need to do any configuration for Spark
 - No, you do not. Because if you indicate to use YARN to run spark, it will automatically use the environment parameter “HADOOP_CONF_DIR” to find the Hadoop configuration files.